



Secret of Eventual Consistency in Apache RocketMQ, with No Budget





Safe Harbor Statement

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Alibaba's products remains at the sole discretion of Alibaba.





Du Heng

Contributor of RocketMQ

Interested in distributed system, such as MQ, Microservice, and has rich experience in performance tuning.





Outline

1. Challenge
2. Overview
3. Architecture
4. Best Practice
5. Future

Aliware阿里中间件







1. Challenge



Challenge



User Case

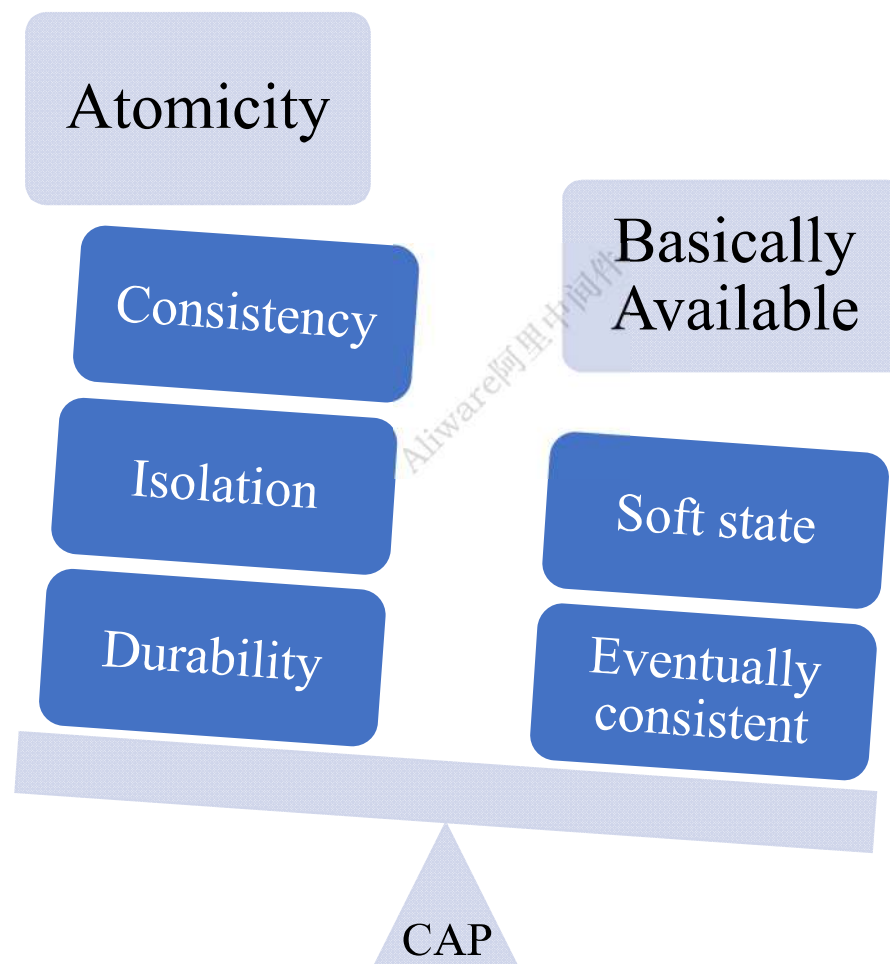




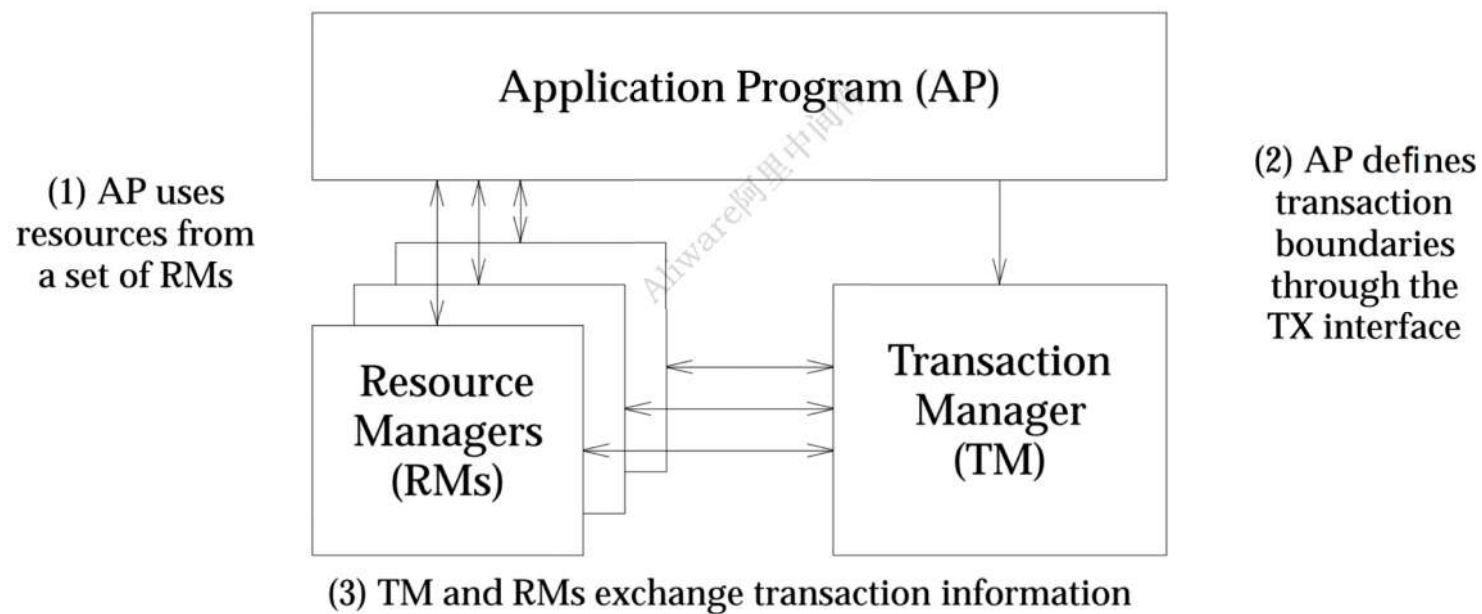
2. Overview




From ACID to BASE



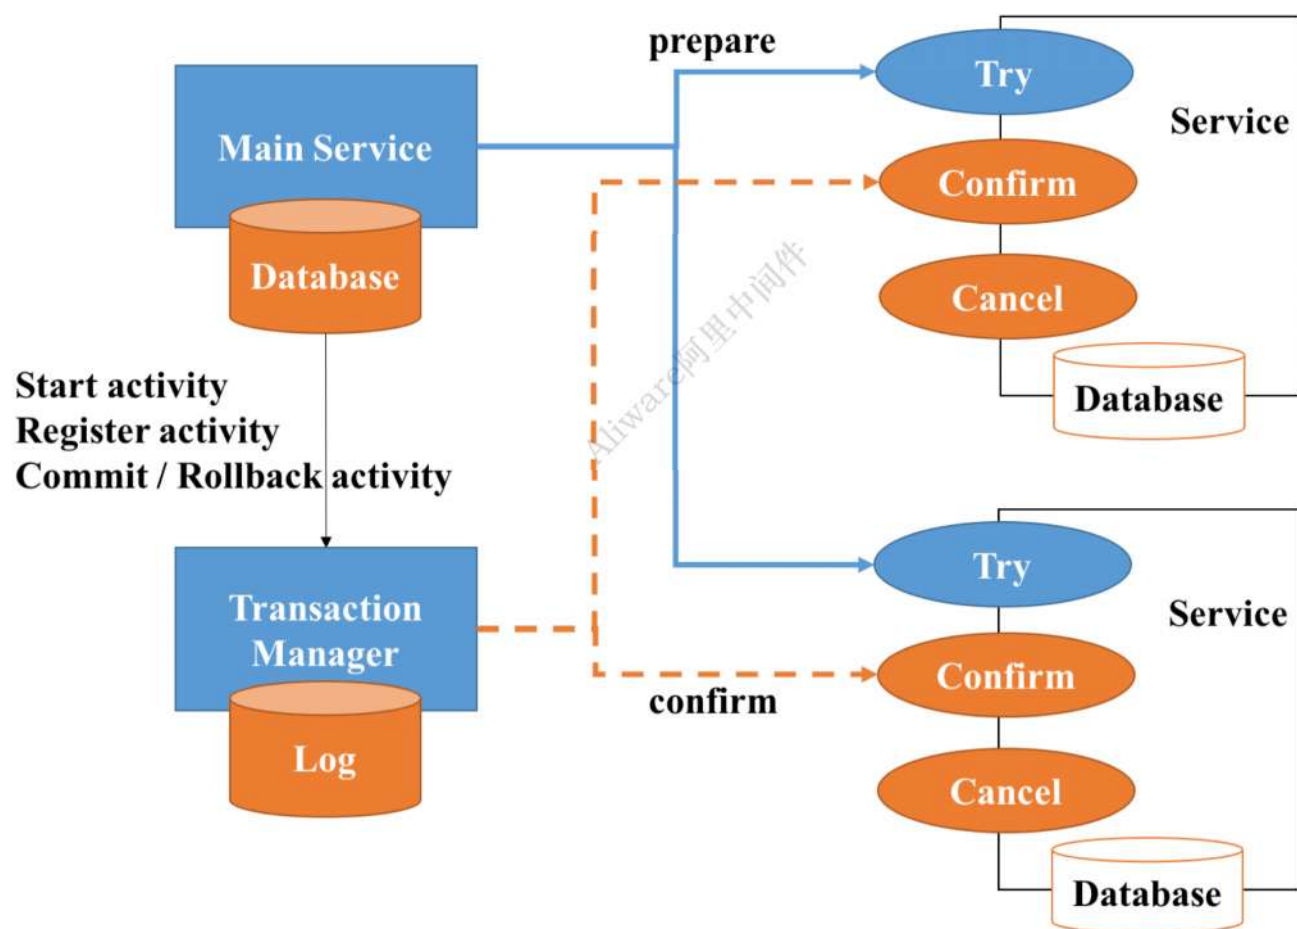
XA



SAGA



TCC

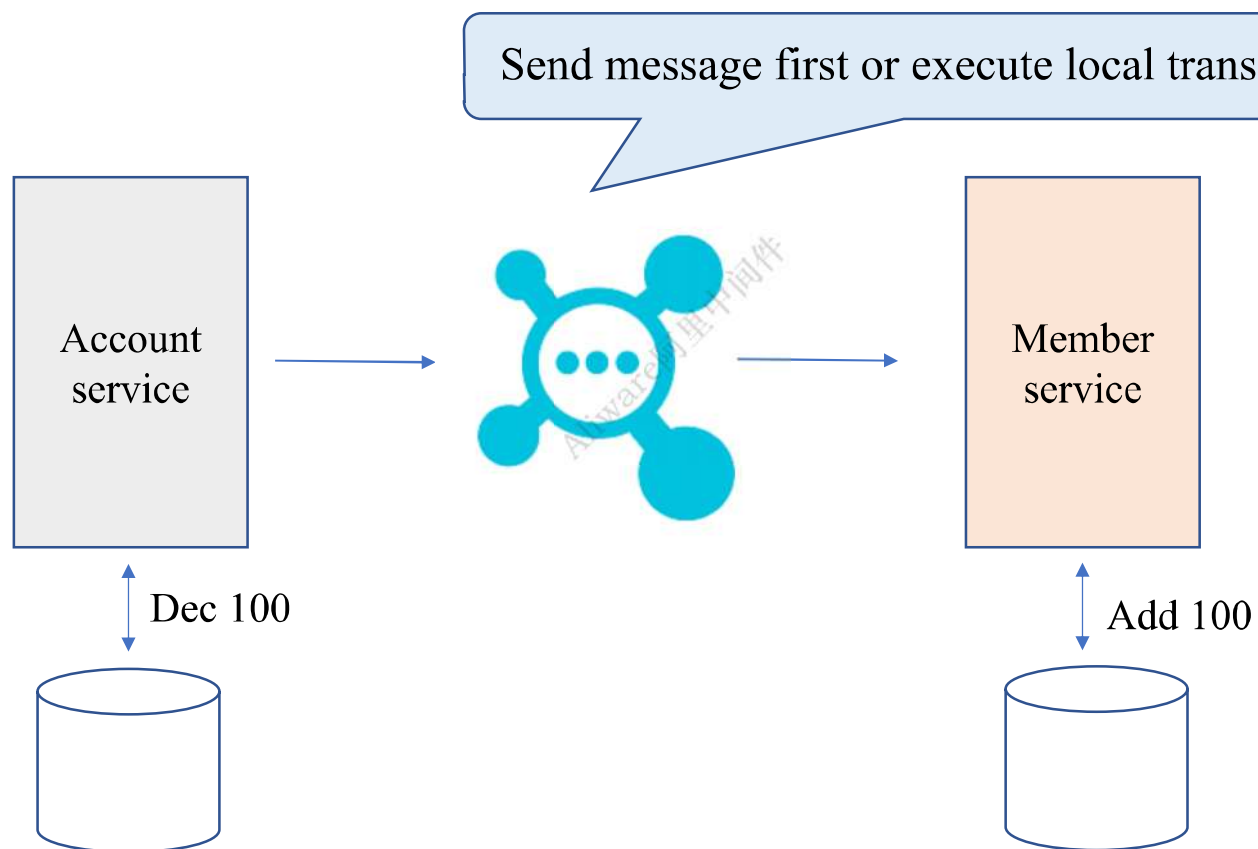




3. Architecture



Transactional Message



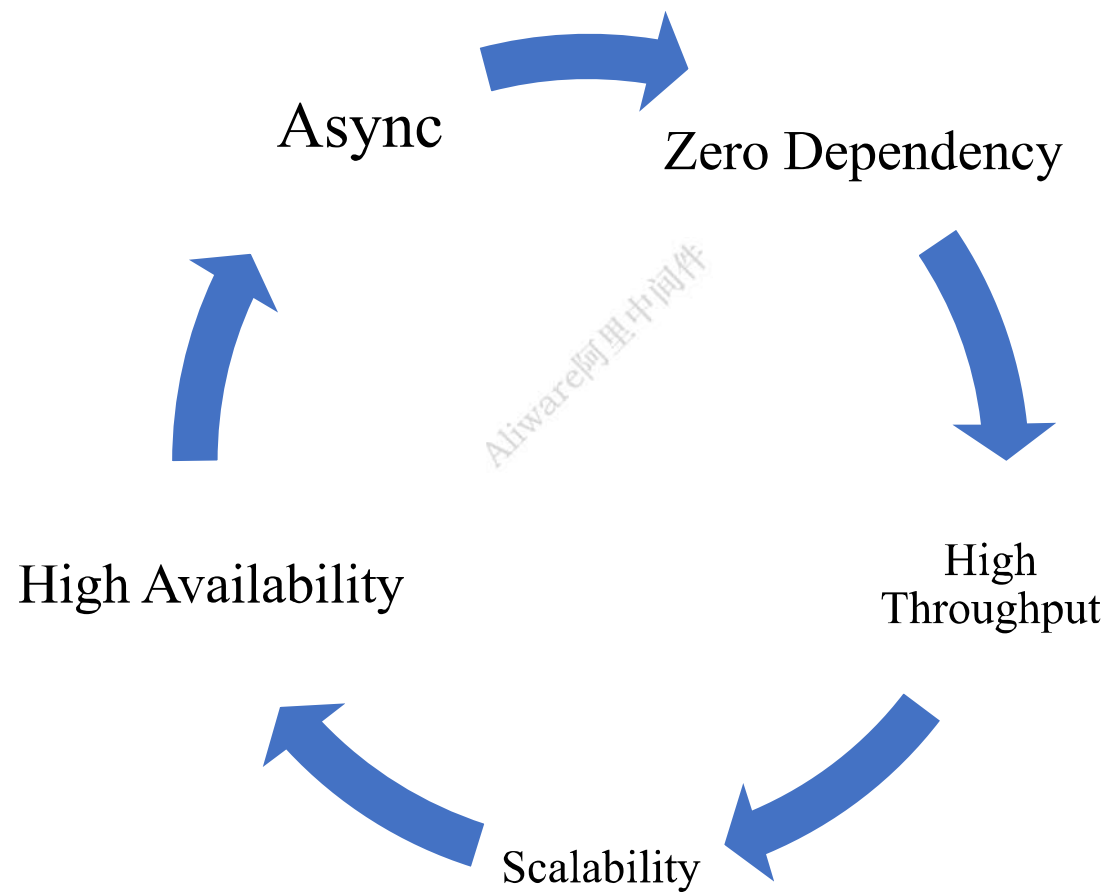


Counterpart Implementation

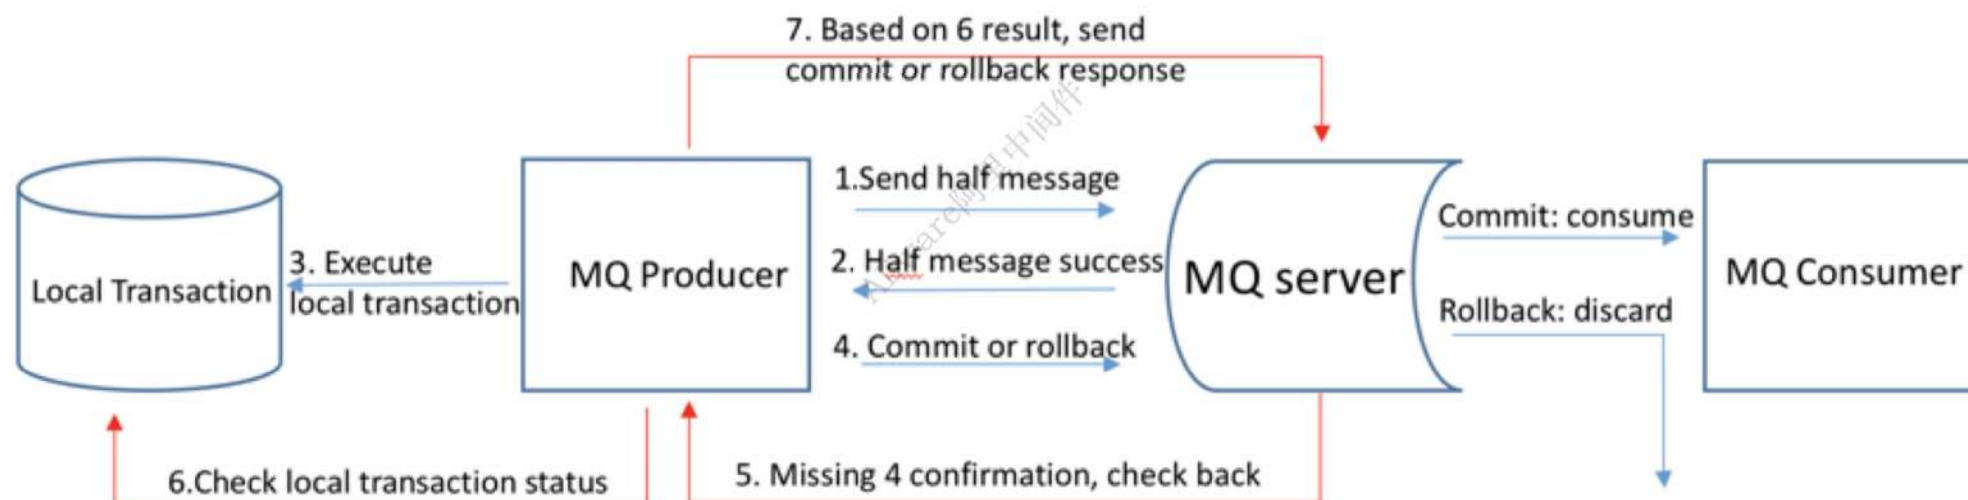
Aliware阿里中间件



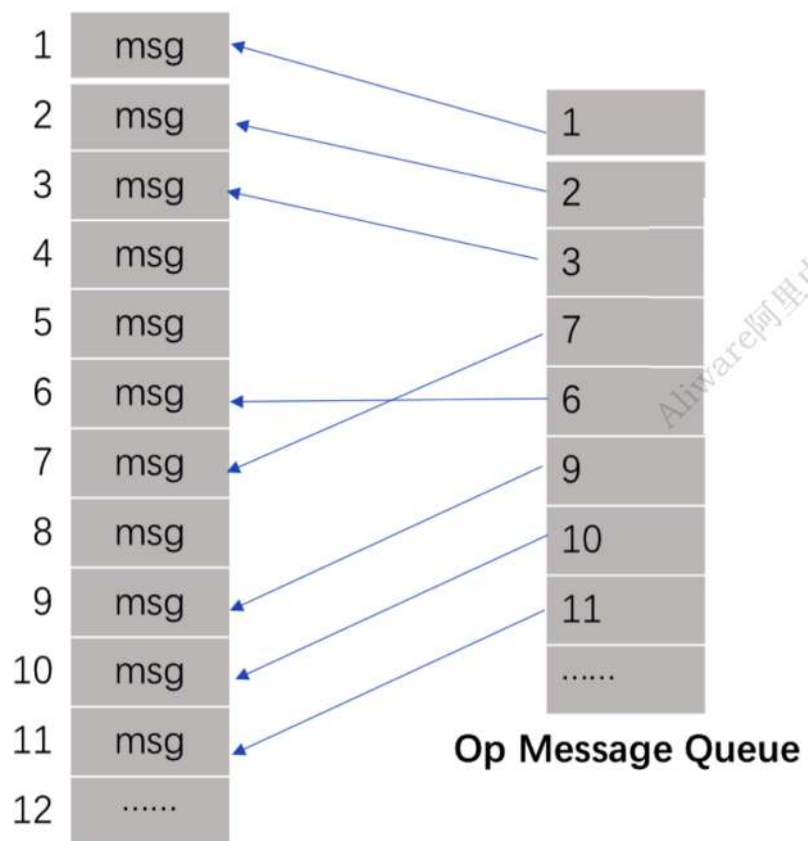
Goal



Procedure



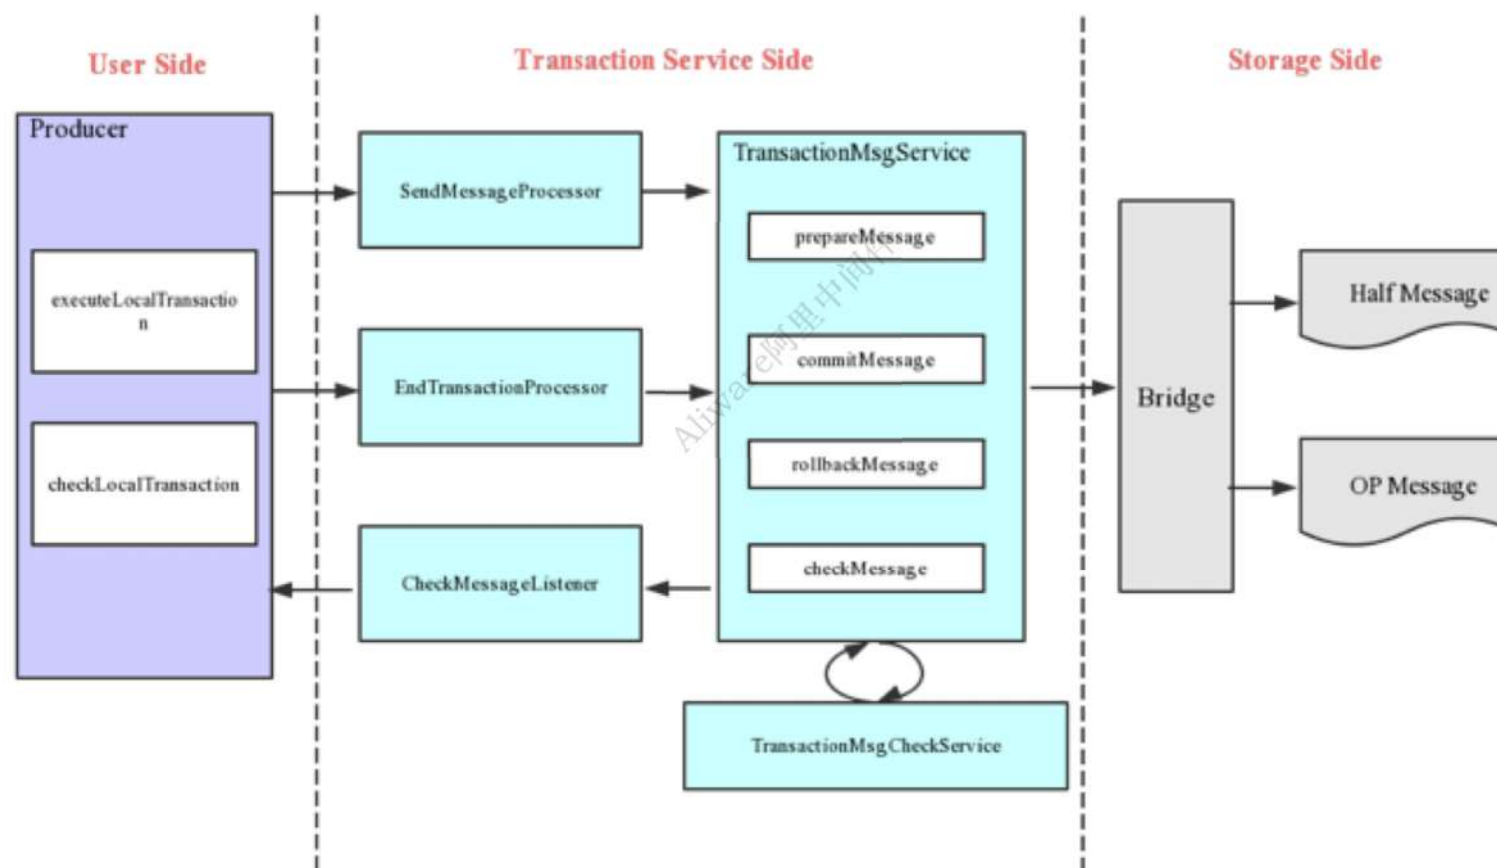
Design



Half Topic: prepare message

Operation Topic: commit/rollback message

Implementation






4. Best Practice



Scenario



Producer Practice

- Use **TransactionalProducer** class to create transactional producer
- Use a specific **producerGroup** name.

```
TransactionListener transactionListener = new TransactionListenerImpl();
TransactionMQProducer producer = new TransactionMQProducer("please_rename_unique_group_name");
ExecutorService executorService = new ThreadPoolExecutor(2, 5, 100, TimeUnit.SECONDS, new ArrayBlockingQueue<Runnable>(2000),
    new ThreadFactory() {
        @Override
        public Thread newThread(Runnable r) {
            Thread thread = new Thread(r);
            thread.setName("client-transaction-msg-check-thread");
            return thread;
        }
    });

producer.setExecutorService(executorService);
producer.setTransactionListener(transactionListener);
producer.start();
```

Producer Practice

- Set the appropriate **thread pool** for your local transaction execution.
- transactionTimeout vs **CHECK_IMMUNITY_TIME_IN_SECONDS**

```
for (int i = 0; i < 10; i++) {  
    try {  
        Message msg =  
            new Message("TopicTest1234", tags[i % tags.length], "KEY" + i,  
                ("Hello RocketMQ " + i).getBytes(RemotingHelper.DEFAULT_CHARSET));  
        msg.putUserProperty(MessageConst.PROPERTY_CHECK_IMMUNITY_TIME_IN_SECONDS, "5000");  
        SendResult sendResult = producer.sendMessageInTransaction(msg, null);  
        System.out.printf("%s%n", sendResult);  
  
        Thread.sleep(10);  
    } catch (MQClientException | UnsupportedEncodingException e) {  
        e.printStackTrace();  
    }  
}
```



Producer Practice

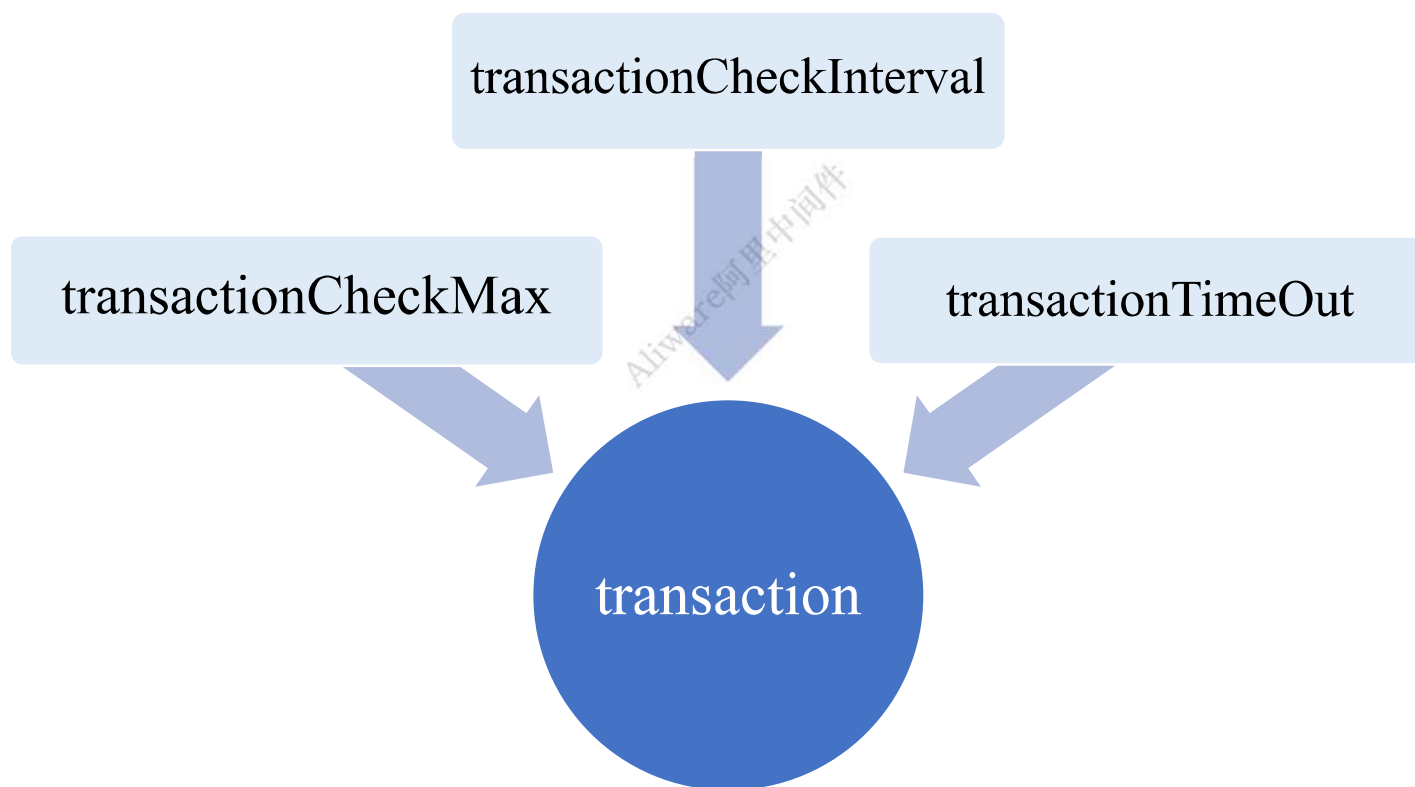
```
public class TransactionListenerImpl implements TransactionListener {  
    private AtomicInteger transactionIndex = new AtomicInteger(0);  
  
    private ConcurrentHashMap<String, Integer> localTrans = new ConcurrentHashMap<>();  
  
    @Override  
    public LocalTransactionState executeLocalTransaction(Message msg, Object arg) {  
        int value = transactionIndex.getAndIncrement();  
        int status = value % 3;  
        localTrans.put(msg.getTransactionId(), status);  
        return LocalTransactionState.UNKNOWN;  
    }  
  
    @Override  
    public LocalTransactionState checkLocalTransaction(MessageExt msg) {  
        Integer status = localTrans.get(msg.getTransactionId());  
        if (null != status) {  
            switch (status) {  
                case 0:  
                    return LocalTransactionState.UNKNOWN;  
                case 1:  
                    return LocalTransactionState.COMMIT_MESSAGE;  
                case 2:  
                    return LocalTransactionState.ROLLBACK_MESSAGE;  
                default:  
                    return LocalTransactionState.UNKNOWN;  
            }  
        }  
        return LocalTransactionState.COMMIT_MESSAGE;  
    }  
}
```

2018/8/2

- Implement **TransactionListener** interface.
 - LocalTransactionState.UNKNOWN
 - LocalTransactionState.COMMIT_MESSAGE
 - LocalTransactionState.ROLLBACK_MESSAGE
- Idempotency (transactionId, message body)

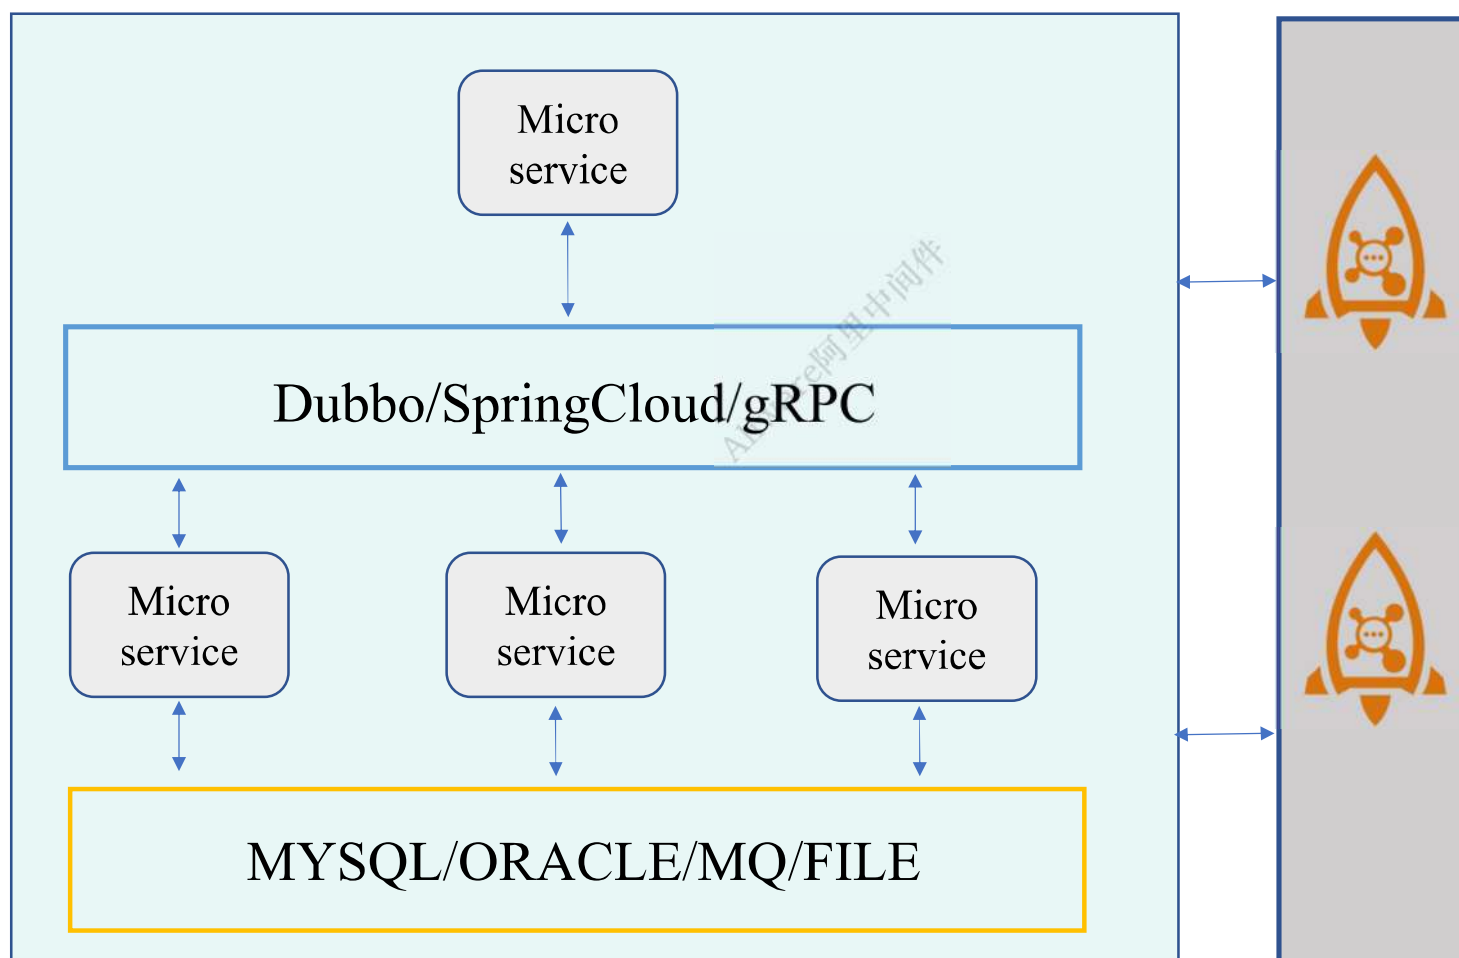


Broker Practice



5. Future

One-stop Distributed Transaction Solution



Thanks !



推荐阅读:

- From Alibaba to Apache: RocketMQ's Past, Present, and Future
- Apache RocketMQ 顶级项目之路
- Apache RocketMQ 背后的设计思路与最佳实践.
- 专访RocketMQ联合创始人: 项目思路、技术细节和未来规划
- 万亿级数据洪峰下的分布式消息引擎
- RocketMQ联合创始人: 选择MQ时, 要注意的有哪些