Pivotal

How to Properly Blame Things for Causing Latency

An Introduction to Distributed Tracing and Zipkin

@Adrian Cole works at Pivotal works on Zipkin

Introduction

introduction
understanding latency
distributed tracing
zipkin
demo
wrapping up



spring cloud at pivotal focused on distributed tracing helped open zipkin

Distributed Tracing

introduction
distributed tracing
zipkin
demo
wrapping up

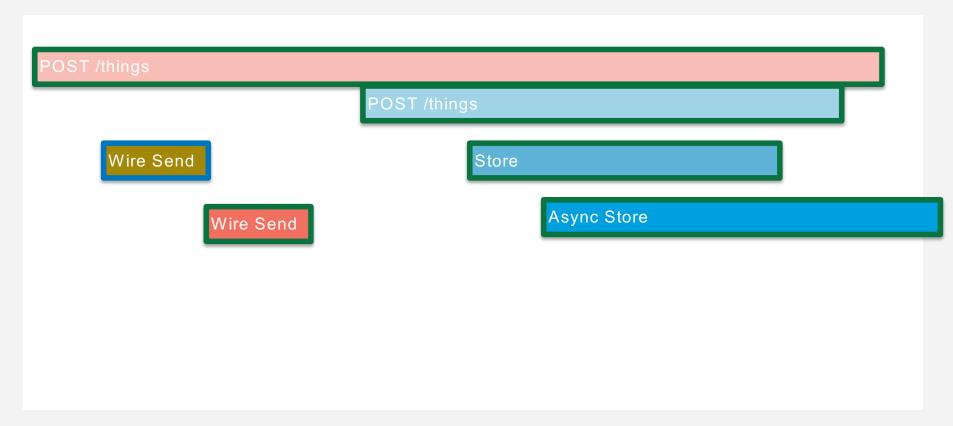
What is Distributed Tracing?

Distributed tracing tracks production requests as they touch different parts of your architecture.

Requests have a unique trace ID, which you can use to lookup a trace diagram, or log entries related to it.

Causal diagrams are easier to understand than scrolling through logs.

Example Trace Diagram



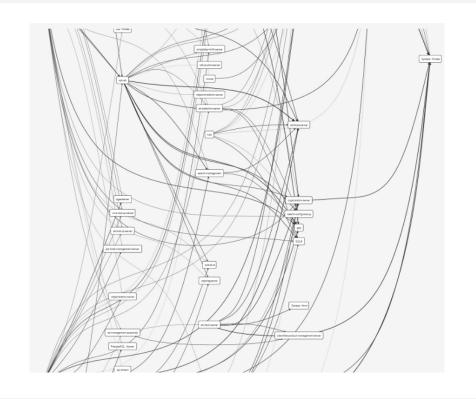
Why do I care?

- Reduce time in triage by contextualizing errors and delays
- Visualize latency like time in my service vs waiting for other services
- Understand complex applications like async code or microservices
- See your architecture with live dependency diagrams built from traces

Example Service Diagram

A tracing system can draw your service dependencies!

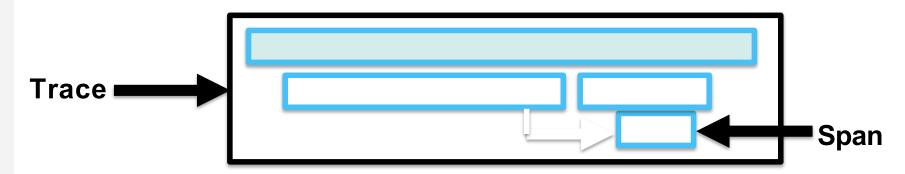
It might resemble your favorite noodle dish!



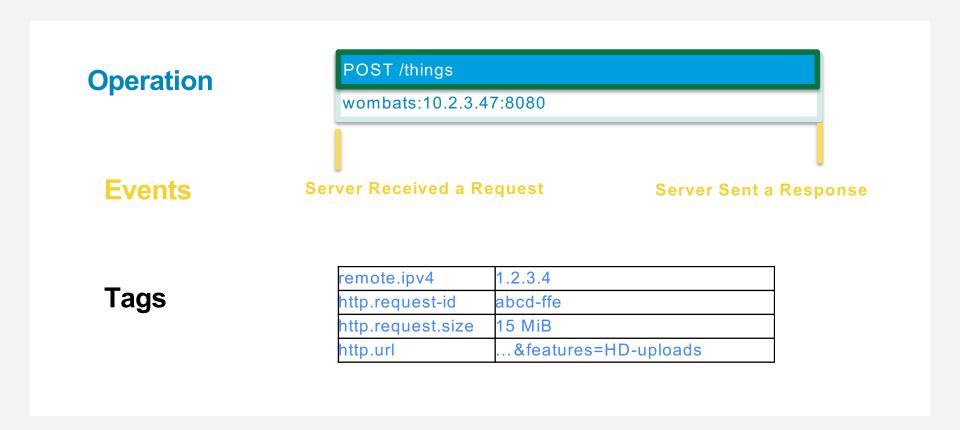
Distributed Tracing Vocabulary

A **Span** is primarily the duration of an operation.

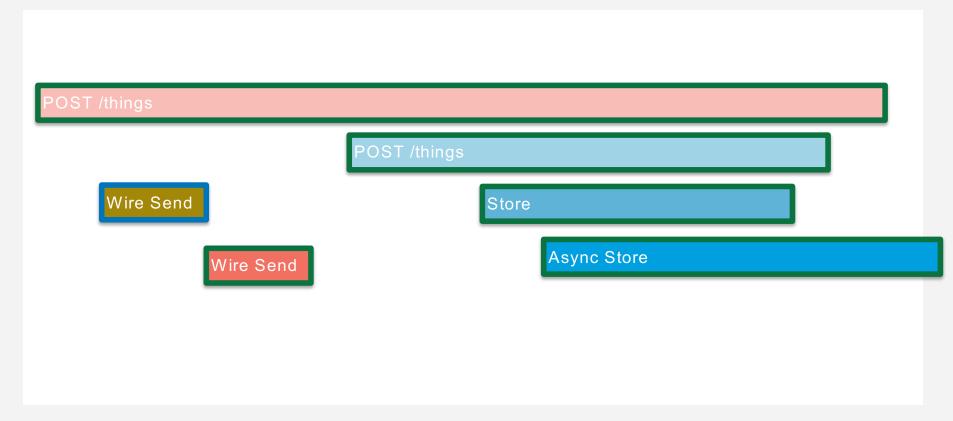
A **Trace** links all spans in a request together by cause.



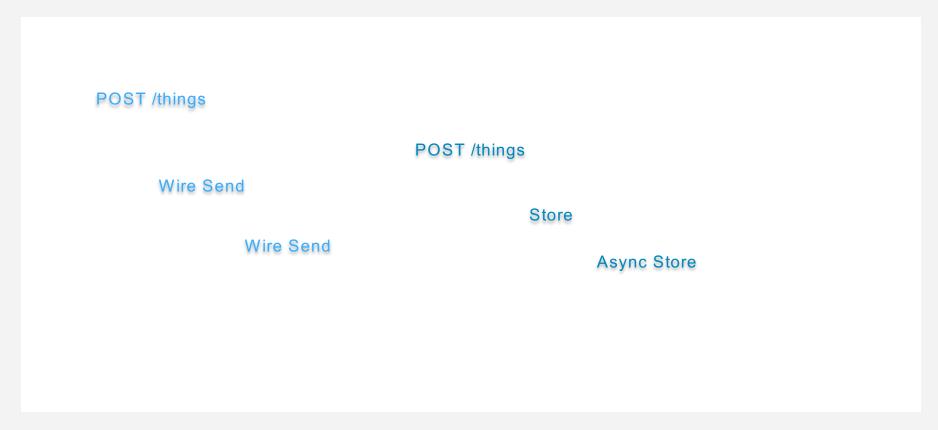
A Span is an individual operation



Trace shows each operation the request caused

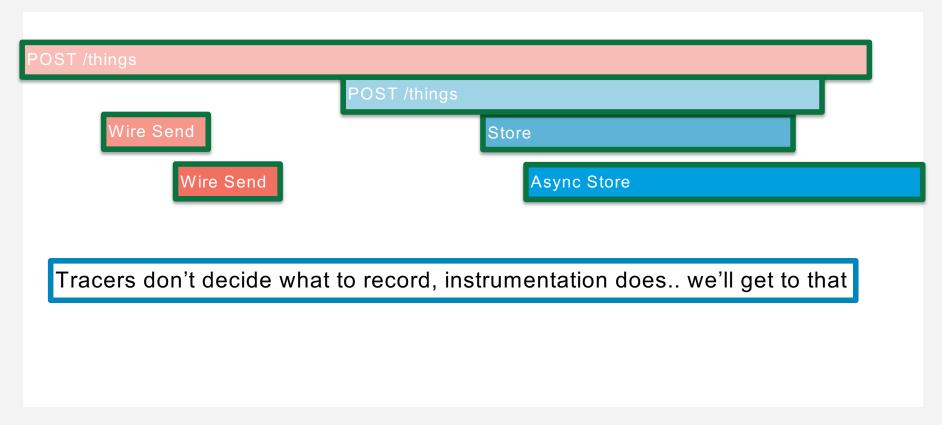


Tracing is capturing important events



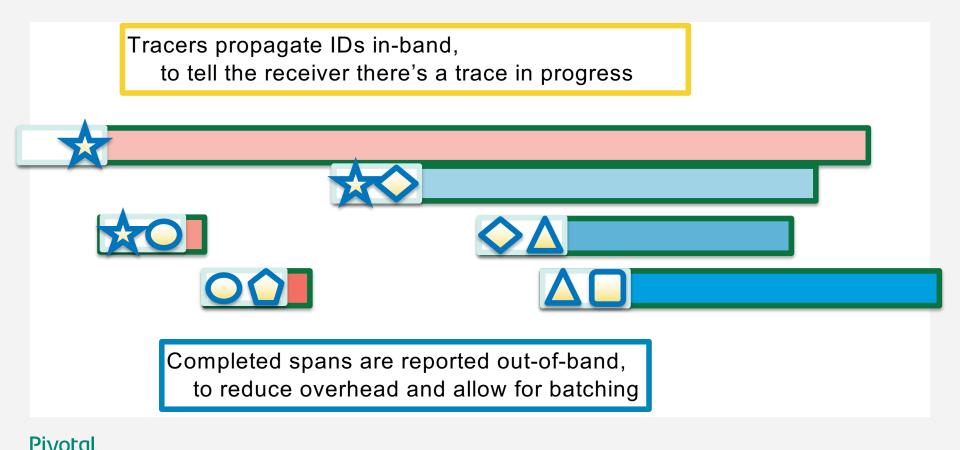
Pivotal.

Tracers record time, duration and host



Pivotal

Tracers send trace data out of process



Tracer vs Instrumentation

A tracer is a utility library similar to metrics or logging libraries.

Instrumentation is framework code that uses a tracer to collect details such as the http url and request timing.

Instrumentation is usually invisible to users

Instrumentation decides what to record

Instrumentation decides how to propagate state

Http Client User Code Trace Instrumentation Zipkin Collector GET /foo record tags add trace headers record timestamp GET /foo X-B3-TraceId: aa X-B3-SpanId: 6b invoke request 200 OK record duration 200 OK asynchronously report span "timestamp": 1483945573944000, "duration": 386000, "annotations": --snip--

Pivotal.

Zipkin

introduction
distributed tracing
zipkin
demo
wrapping up

Zipkin is a distributed tracing system

tornado-server

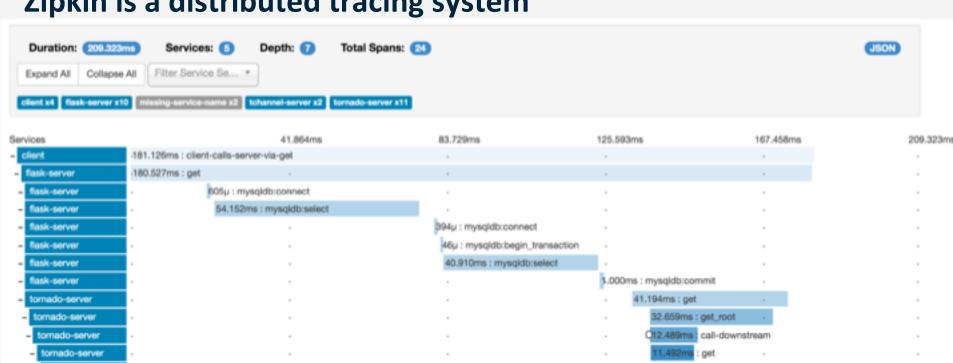
tomado-server

tornado-server

tornado-server tornado-server

tornado-server

tchannel-server



105µ: tomado-x2

get 85µ:tornado-x3 -

C29.816ms : call-tchannel-

call-downstream

C12.153ms : call_in_request_context

9.712ms: endpoint

O11,494ms

10.511ms

Zipkin lives in GitHub

Zipkin was created by Twitter in 2012 based on the Google Dapper paper. In 2015, OpenZipkin became the primary fork.

OpenZipkin is an org on GitHub. It contains tracers, OpenApi spec, service components and docker images.

https://aithub.com/openzipkin

Zipkin Architecture

Tracers **report** spans HTTP or Kafka.

Servers **collect** spans, storing them in MySQL, Cassandra, or Elasticsearch.

Users **query** for traces via Zipkin's Web UI or Api.

Amazon Azure Docker Goodle Kubernetes Mesos Spark

Instrumented client (Reporter) Instrumented server Non-instrumented server (Reporter) Zipkin Transport UI Collector API Storage Database

Pivotal

Zipkin has starter architecture

Tracing is new for a lot of folks.

For many, the MySQL option is a good start, as it is familiar.

```
services:
    storage:
    image: openzipkin/zipkin-mysql
    container_name: mysql
    ports:
        - 3306:3306
    server:
        image: openzipkin/zipkin
        environment:
        - STORAGE_TYPE=mysql
        - MYSQL_HOST=mysql
        ports:
        - 9411:9411
        depends_on:
        - Storage
```

Zipkin can be as simple as a single file

```
$ curl -SL 'https://search.maven.org/remote_content?g=io.zipkin.java&a=zipkin-server&v=LATEST&c=exec' > zipkin.jar
$ SELF_TRACING_ENABLED=true java -jar zipkin.jar
                                   ****
                                   ****
  *****************************
       ****
                                   ****
                                    **
            *****
                                  (v1.5.4.RELEASE)
:: Powered by Spring Boot ::
2016-08-01 18:50:07.098 INFO 8526 --- [ main] zipkin.server.ZipkinServer example/zipkin.jar started by acole in /Users/acole/oss/sleuth-webmvc-example)
                                                                                          : Starting ZipkinServer on acole with PID 8526 (/Users/acole/oss/sleuth-webmvc-
-snip-
                                                                      $ curl -s localhost:9411/api/v2/services|jq .
                                                                          "gateway"
```

Brave: the most popular Zipkin Java tracer

- Brave OpenZipkin's java library and instrumentation
 - Layers under projects like Armeria, Dropwizard, Play
- Spring Cloud Sleuth automatic tracing for Spring Boot
 - Includes many common spring integrations
 - Starting in version 2, Sleuth is a layer over Brave!

```
c, c#, erlang, javascript, go, php, python, ruby, too
```

Some notable open source tracing libraries

- OpenCensus Observability SDK (metrics, tracing, tags)
 - Most notably, gRPC's tracing library
 - Includes exporters in Zipkin format and B3 propagation format
- OpenTracing trace instrumentation library api definitions
 - Bridge to Zipkin tracers available in Java, Go and PHP
- SkyWalking APM with a java agent developed in China
 - Work in progress to send trace data to zipkin
- Kamon AkKa Monitoring: trace and metrics specializing in scala
 - Uses B3 propagation and has a Zipkin export plugin





Pivotal

Transforming How The World Builds Software